# Long-Context QA with Modern Hopfield Network

1st Niranjan Ambi
*Department of Computer Science*
*KIT College of Engineering, Kolhapur*
Kolhapur, India
niranjan@example.com

2nd Tirth Katkar
*Department of Computer Science*
*KIT College of Engineering, Kolhapur*
Kolhapur, India
tirth@example.com

3rd Pratham Kadam
*Department of Computer Science*
*KIT College of Engineering, Kolhapur*
Kolhapur, India
pratham@example.com

4th Shreya Gunjkar
*Department of Computer Science*
*KIT College of Engineering, Kolhapur*
Kolhapur, India
shreya@example.com

***Abstract*—Open-domain question answering (QA) systems based on standard autoregressive architectures face computational bottlenecks due to the quadratic complexity of self-attention. Traditional Retrieval-Augmented Generation (RAG) mitigates these limits using vector embeddings but assumes conditional independence between retrieved chunks. This assumption can lead to redundant retrieval and poor reasoning in multi-hop scenarios. This paper explores the integration of Modern Continuous Hopfield Networks as an energy-based associative memory retriever. By using a LogSumExp energy function, this architecture achieves exponential storage capacity and mathematically mirrors Transformer attention, allowing systems to converge on stable combinations of stored facts. We review the theoretical evolution from classical binary networks to continuous models, outline PyTorch implementation strategies, address practical limitations like metastability and rank collapse, and evaluate performance on multiple instance learning benchmarks like SQuAD.**

***Index Terms*—Question Answering, Modern Hopfield Networks, Associative Memory, Retrieval-Augmented Generation, Natural Language Processing, Transformer Architecture**

## I. The Need for Scalable Long-Context Retrieval

Natural language processing models must often process and retrieve information across large textual contexts. As systems shift toward comprehensive, open-domain question answering (QA), the limitations of standard autoregressive architectures become clear. Foundational models built on the standard Transformer architecture face computational and memory bottlenecks when exposed to massive text corpora. A primary limitation is the self-attention mechanism, which scales quadratically with the length of the input sequence. Deploying these models for long-context QA across thousands of documents requires external memory augmentations to bypass strict token limits.

Retrieval-Augmented Generation (RAG) is a common architectural approach for bridging this gap. Traditional RAG frameworks circumvent sequence length limitations using dense vector embeddings. In this approach, massive textual corpora are segmented into chunks, embedded into a continuous vector space using a dual-encoder architecture like the Dense Passage Retriever, and stored in a vector database. When a query is issued, it is embedded, and the system executes a Maximum Inner Product Search (MIPS) or calculates cosine similarities to retrieve the top-$k$ most relevant document chunks. These chunks are concatenated into a prompt and passed to the generative language model to synthesize an answer.

While RAG with vector embeddings is widely adopted due to its modularity and low latency, it has theoretical limitations in reasoning tasks. Traditional retrievers operate under an assumption of conditional independence. Each document chunk is evaluated against the query vector in isolation, independent of other retrieved chunks.

In complex QA scenarios, such as those requiring multi-hop reasoning, temporal synthesis, or the aggregation of distributed facts, this independent scoring mechanism often fails. It retrieves redundant information that answers only a fraction of the query while missing correlated pieces of knowledge necessary for a complete response. Additionally, when relevant information is retrieved, models often exhibit the "lost in the middle" phenomenon, where the attention mechanism fails to adequately weight facts positioned in the center of the concatenated context.

Associative Memory Retrievers using Modern Hopfield Networks (MHNs) offer an alternative solution. Unlike standard RAG methodologies that rely on static distance metrics, MHNs introduce dynamic, energy-based recall mechanisms. In a Hopfield formulation, the massive textual corpus acts as a dense associative memory. The user's query is treated as an incomplete initial state. Through a recurrent update rule, the network iteratively minimizes a global energy function, allowing the query state to dynamically evolve and gravitate toward a stable attractor. This local energy minimum corresponds to a coherent and relevant combination of stored patterns.

Constructing a QA system utilizing PyTorch-Hopfield layers involves a shift from discrete database lookups to continuous dynamical systems optimization. Unnormalized energy models present specific training complexities, often requiring Monte Carlo sampling to approximate partition functions. However, the integration of Modern Hopfield Networks provides a math-

ematically grounded framework capable of handling memory-intensive tasks. This paper examines the theoretical foundations, implementation strategies, benchmarking methodologies on datasets like SQuAD, and the empirical frameworks required to architect an Associative Memory Retriever for long-context QA.

## II. Theoretical Foundations: From Binary States to Exponential Continuous Memory

This section traces the development of associative memory models and how their computational capacity limitations were addressed over time.

### A. The Classical Binary Hopfield Network

The classical Hopfield network, conceptualized in 1982, was introduced as a recurrent artificial neural network that mathematically modeled associative, content-addressable memory. The network consisted of a single layer of discrete neurons, where each neuron functioned as a binary threshold unit, taking on states of either $+1$ or $-1$.

The architecture was defined by bidirectional and symmetric connections ($w_{ij} = w_{ji}$) with no self-connections ($w_{ii} = 0$). Patterns were associatively learned and stored in the network utilizing a Hebbian learning algorithm, which computed the outer product of the pattern vectors to construct the global weight matrix $W$.

The network is defined by its energy landscape. For a given state vector $\xi$, the network iteratively updated individual neuron states to minimize the overall energy, governed by the update rule $\xi^{(t+1)} = \text{sgn}(W\xi^{(t)} - b)$. The network dynamics guaranteed convergence to local energy minima, corresponding to the stored memory patterns. This allowed the system to recover complete memories from partial or noisy inputs.

However, the classical binary Hopfield network had a restricted storage capacity. The number of memory patterns $C$ that the network could store without error scaled linearly with the number of dimensions or neurons $N$, formally established at $C \approx 0.138N$. Exceeding this threshold caused the defined attractor basins to overlap, creating spurious local minima. The network would suffer from interference, returning mixed outputs rather than discrete memories. This linear scaling law limits the classical network's utility in modern NLP tasks that require storing millions of high-dimensional textual embeddings.

### B. Dense Associative Memories and Polynomial Interactions

The transition to Dense Associative Memories occurred in 2016 when Krotov and Hopfield demonstrated that the linear capacity barrier could be broken by altering the energy function. Instead of relying on pairwise, quadratic interactions between neurons, they introduced higher-order polynomial interaction functions.

By defining the energy of the state relative to the stored patterns using polynomial activation functions of degree $n$, the storage capacity of the network transitioned into the super-linear regime. While this formulation proved that associative memories could theoretically store more patterns than neurons, the reliance on arbitrary polynomial degrees made the architecture computationally rigid.

### C. Modern Continuous Hopfield Networks

A major development for applying associative memory to long-context QA was the introduction of Modern Continuous Hopfield Networks by Ramsauer et al. in 2020. This framework transitioned the network from binary variables to continuous states, allowing it to interface directly with the high-dimensional continuous representations produced by language models.

Ramsauer et al. replaced the polynomial interaction terms with a non-linear exponential function, specifically the Log-SumExp (LSE) function. Let $X = (x_1, x_2, \dots, x_N)$ represent the matrix of $N$ stored memory patterns, where each pattern $x_i \in \mathbb{R}^d$. Let $\xi \in \mathbb{R}^d$ represent the current continuous state (the query vector). The energy function $E$ of the Modern Hopfield Network is defined as:

$$E = -\beta^{-1} \log \left( \sum_{i=1}^{N} \exp(\beta x_i^T \xi) \right) + \frac{1}{2}\xi^T\xi + \beta^{-1} \log N + \frac{1}{2}M^2 \tag{1}$$

Here, $\beta$ represents an inverse temperature parameter that scales the dot product similarity, and $M$ is defined as the largest norm of any stored pattern ($M = \max_i ||x_i||$). The term $\beta^{-1}\log(\dots)$ acts as a smooth maximum function, exponentially weighting the stored patterns that exhibit the highest inner product with the current state $\xi$. The inclusion of the quadratic regularizer $\frac{1}{2}\xi^T\xi$ ensures that the energy function is globally bounded and that the update dynamics will converge to a stable fixed point on the energy landscape.

Mathematical analysis proves that for certain choices of activation functions, the storage capacity of the Modern Hopfield Network becomes exponential relative to the feature dimension $d$. The capacity scales proportionately to $2^{d/2}$. This exponential capacity allows the network to store and retrieve millions of factual text embeddings without interference.

### D. Isomorphism of Hopfield Dynamics and Transformer Attention

The viability of an Associative Memory Retriever rests on a mathematical equivalence discovered within the continuous formulation. When calculating the derivative of the Modern Hopfield energy function with respect to the state vector $\xi$ to derive the update rule, the resulting equation is:

$$\xi^{(t+1)} = X\text{softmax}(\beta X^T \xi^{(t)}) \tag{2}$$

This update rule is mathematically isomorphic to the self-attention mechanism that powers the Transformer architecture. In a standard Transformer, the attention mechanism computes output sequences by multiplying a value matrix $V$ by the softmax of the dot product of a query matrix $Q$ and a key matrix $K$, scaled by a dimension factor. When mapped to the Hopfield update rule, the query $Q$ corresponds to the initial

state $\xi^{(t)}$, while the key $K$ and value $V$ matrices correspond identically to the stored memory patterns $X$. The inverse temperature $\beta$ corresponds to the scaling factor $\frac{1}{\sqrt{d_k}}$.

This equivalence reveals that the attention heads within modern LLMs execute a single update step of a Modern Continuous Hopfield Network. Recognizing this allows developers to decouple the memory retrieval process from the generative feed-forward pipeline. Instead of forcing a Transformer to hold massive contexts within its internal self-attention layers, external PyTorch-Hopfield layers can be deployed for dense associative retrieval. These layers function as a plug-in replacement for standard attention layers and feed the energy-stabilized factual representations back into the LLM decoder.

## III. Energy-Based Recall vs. Standard Retrieval-Augmented Generation

In standard RAG, the retrieval phase is a deterministic geometric lookup. The query is embedded as a vector $q$, and the corpus chunks are embedded as $k_i$. The retriever calculates the dot product $q^T k_i$ for all $i$, ranks them, and returns the top-$k$. This treats every chunk as conditionally independent.

If a user poses a complex question requiring multi-hop reasoning, a standard RAG retriever searching for vector similarity against the raw query will heavily weight chunks containing matching keywords. It may retrieve redundant chunks while failing to retrieve the specific text detailing the necessary logical connections if those chunks have low lexical overlap with the initial query.

Energy-based retrievers eliminate this conditional independence by modeling the candidate retrieval results as a cohesive ensemble. In a Hopfield formulation, the initial query is the unstable state $\xi^{(0)}$. When the update rule $\xi^{(1)} = X\text{softmax}(\beta X^T \xi^{(0)})$ is applied, the new state becomes a linear combination of all stored memories, heavily weighted toward those in the immediate attractor basin.

As the state iteratively updates, it absorbs the features of the activated memories. When $\xi^{(2)}$ is computed, the query vector has changed and actively searches the memory matrix $X$ using the newly acquired contextual features.

This iterative descent into an energy minimum allows the network to operate in a metastable state. This is an energy minimum that averages over a logically correlated subset of patterns rather than collapsing into a single isolated memory. The retrieved context forms an energetically stable consensus of interlinked facts. Consequently, this multi-hop alignment reduces hallucinations, as the final representation is grounded in a network of continuous patterns rather than a list of independently scored sentences.

## IV. Architecting the PyTorch-Hopfield QA System

Translating theoretical attractor dynamics into an operational deep learning architecture involves leveraging optimized linear algebra operations to perform the updates required by the network.

### A. Encoding the Corpus and Defining the Memory Matrix

The massive textual corpus is preprocessed into discrete semantic units, such as cohesive sentences or localized paragraphs. A pre-trained dense encoder embeds these units into high-dimensional vectors. These vectors are concatenated to form the static memory matrix $X \in \mathbb{R}^{d \times N}$, where $d$ is the hidden embedding dimension and $N$ is the total number of text chunks. The memory matrix $X$ resides outside the generative computational graph, allowing $N$ to scale effectively.

### B. Query Initialization and Energy Minimization

When a user submits a question, it is processed by the encoder to generate the initial state vector, denoted as state ($\xi \in \mathbb{R}^d$). This tensor is passed into the PyTorch-Hopfield layer, which executes the forward pass utilizing the LSE-based energy calculations:

```python
import torch
import torch.nn as nn

class Hopfield(nn.Module):
    def __init__(self, X):
        super(Hopfield, self).__init__()
        self.X = X  # Shape: (d, N)

    def forward(self, state):
        # xi_new = X * softmax(X^T * xi)
        attn = torch.softmax(self.X.T @ state, 0)
        new_state = self.X @ attn
        return new_state
```

Driven by the LogSumExp function, convergence to a stable fixed point typically occurs in a single update step, maintaining $O(N \times d)$ computational complexity per query.

### C. Extractive Decoding for Question Answering

Once the Hopfield layer has stabilized the query state into a contextually enriched representation, this tensor is mapped to a natural language answer. For extractive QA paradigms, the new state is concatenated with the raw embeddings of the most highly activated memory chunks.

This combined representation is passed into a lightweight processing head, such as a multi-layer perceptron (MLP) or a bidirectional LSTM. The head projects the representations through a linear classification layer, applying a softmax function to generate start and end probability distributions over the tokens to extract the specific factual span. Because the Hopfield layer handles the multi-hop contextual alignment, the extraction head requires fewer parameters than a standard Transformer block to achieve accuracy.

## V. Benchmarking on SQuAD: Formulating Memory as Multiple Instance Learning

The Stanford Question Answering Dataset (SQuAD) serves as a benchmark for evaluating reading comprehension and factual retrieval. Evaluating a long-context QA system on SQuAD requires formulating the task properly.

Standard architectures evaluate SQuAD by concatenating the query and the context into a single string. This strategy fails when the context expands to document-level texts due to

token limit truncation. To bypass this, researchers formulate the long-context SQuAD task as a Multiple Instance Learning (MIL) problem. In the MIL paradigm, a massive document is treated as a "bag" containing sentences as instances. The system must predict a label for the entire bag, but the specific instance containing the fact is hidden among irrelevant background instances.

Standard models often struggle at massive MIL tasks due to the low witness rate. However, the Modern Hopfield Network isolates the high-energy witness pattern while ignoring the irrelevant instances.

When benchmarked against SQuAD using Exact Match (EM) and F1 Score metrics, baseline models like standard Bidirectional LSTMs struggle to surpass 55% F1. Advanced Multi-Perspective Matching architectures plateau around 68% F1 on the training data. The integration of Hopfield layers eliminates token truncation errors and utilizes energy minimization to resolve noisy sentences, improving upon these baselines in both F1 and Exact Match metrics across massive documents.

## VI. Advanced Architectures: DeepRC, Energy Transformer, and AiT

The translation of continuous Hopfield theories into architectures has yielded frameworks designed to tackle the limitations of self-attention in QA and sequence modeling.

The application of Hopfield networks to massive dataset QA originated from the DeepRC framework. Initially designed for computational biology to classify immune repertoires, the methodology showed that integrating attention with a fixed query could solve tasks with massive instances. DeepRC demonstrated that a PyTorch-based Hopfield layer could identify discriminating sequence motifs among noisy bags, establishing a blueprint for long-context SQuAD retrieval.

Building upon this, the Energy Transformer (ET) modified the feed-forward pipeline. In standard models, attention, MLPs, and layer normalizations are interleaved. The ET replaces this sequence with an architecture designed to minimize a global energy function representing the relationships between tokens. A single ET block recurrently iterates token representations until they converge to a fixed point attractor state. The strict adherence to global energy minimization makes the model structurally rigid, prompting the development of hybrid approaches.

The Associative Transformer (AiT) provides a parameter-efficient hybrid approach. The AiT leverages a learnable explicit memory comprising specialized priors that act as attractors within the continuous Hopfield network. The AiT uses a bottleneck mechanism that forces inputs to compete for information writing, after which the Hopfield layer performs memory-based token reconstruction. This approach integrates the Hopfield network directly into the sparse attention mechanism.

## VII. Navigating Practical Limitations: Sparsity, Rank Collapse, and Compression

Deploying associative memory retrievers over traditional RAG engines presents implementation challenges related to the Hopfield energy landscape.

### A. Resolving Metastability with Sparse Modern Hopfield Networks

The standard softmax probability mapping driven by the LogSumExp function is susceptible to metastability if the input patterns are not well-separated. A pattern is well-separated if its dot product with itself is significantly higher than its dot product with any other pattern.

In textual corpora used for QA, documents are frequently not well-separated. Articles discussing similar historical events share semantic overlap. When a query is introduced, weak pattern separation can cause the network to settle into spurious metastable states, blending facts.

This is addressed by Sparse Modern Hopfield Networks. By substituting the standard softmax mapping with the $\alpha$-entmax family of sparse probability mappings, the energy functional is regularized by entropy. The $\alpha$-entmax formulation forces the network to assign exactly zero weight to irrelevant patterns, enforcing a localized attractor basin that delays metastability and prevents semantic blending.

### B. Mitigating Rank Collapse via Hidden State Dynamics

Deep QA architectures that stack Hopfield layers with standard Transformer blocks can experience rank collapse. This occurs when feature representations rapidly degenerate into a rank-1 matrix as depth increases, leading to token uniformity where query tokens look mathematically identical.

Applying a moving average to update the hidden states across layers explicitly incorporates Hopfield recursion into the attention matrix. Maintaining this continuous hidden-state stream ensures that as the query navigates through layers of energy minimization, its mathematical representation does not homogenize.

### C. Continuous-Time Memory Compression

A physical limitation of the PyTorch-Hopfield implementation is the requirement to store the continuous dataset explicitly in the memory matrix $X$. The physical VRAM allocation scales linearly with the size of the text corpus, which becomes a hardware constraint when scaling to millions of documents.

Compressing large discrete Hopfield memories into continuous-time representations offers a structural optimization. The traditional probability mass function utilized in standard updates is replaced with a probability density function integrated over the continuous memory space. This shift reduces computational memory costs while maintaining the integrity of the attractor dynamics.

### D. Advanced Information Synthesis: Graph Modeling and Unnormalized Training

Deploying an energy-based retriever requires methodologies for handling complex signals. A query might ask for a sequence of events, requiring chronological understanding.

The AssoMem architecture addresses this by mapping retrieved utterances onto an associative memory graph, integrating relevance, importance, and temporal alignment using a Mutual Information fusion strategy. Biasing the energy landscape with temporal signals forces the query state toward attractor basins that are both semantically relevant and chronologically accurate.

However, training unnormalized energy models introduces computational demands. Optimizing the Hopfield energy landscape using Maximum Likelihood Estimate requires the continuous estimation of the partition function. Engineers often employ Monte Carlo sampling methods, specifically Importance Sampling or Metropolis Independent Sampling, to approximate gradients during backpropagation.

This sampling complexity can destabilize training. Using a Joint Stochastic Approximation algorithm within a semi-supervised framework helps mitigate this. The system undergoes supervised pre-training on labeled data and subsequently uses the frozen retrieval parameters to calculate importance weights. While this improves accuracy metrics, the architectural overhead keeps standard RAG embeddings as the default in many general applications.

## VIII. The Future of Episodic Memory Tasks

The architecture of a PyTorch-Hopfield QA system bridges statistical language modeling and energy-based reasoning. While foundational models built on standard Transformers continue to be widely used, limitations regarding attention spans and logic routing persist.

The Associative Memory Retriever processes information recurrently, moving toward an answer through energy minimization. Frameworks integrating Hopfield retrievers alongside Contextual Knowledge Sets enable the dynamic decomposition of queries into reasoning chains, which can aid accuracy and inference latency.

The transition to the continuous capacity of Dense Associative Memories aligns the Hopfield update rule with the Transformer attention mechanism, allowing the integration of external memory layers into existing pipelines. As research refines sparse regularizations and memory compressions, the Associative Memory Retriever is positioned to operate effectively in massive episodic memory tasks.

## References


[1] Entriever: Energy-based Retriever for Knowledge-Grounded Dialog, arXiv, accessed February 18, 2026, https://arxiv.org/abs/2506.00585

[2] Retrieval Meets Long Context Large Language Models, ICLR Proceedings, accessed February 18, 2026.

[3] Modern Hopfield Networks for graph embedding, Frontiers, accessed February 18, 2026.

[4] Hopfield Networks is All You Need, arXiv, accessed February 18, 2026, https://arxiv.org/abs/2008.02217

[5] Hopfield network, Wikipedia, accessed February 18, 2026.

[6] A quantum Hopfield associative memory implemented on an actual quantum processor, PMC, accessed February 18, 2026.

[7] Reasoning behind performance improvement with hopfield networks, AI Stack Exchange, accessed February 18, 2026.

[8] Sepp Hochreiter on Parallels Between Attention Mechanisms and Modern Hopfield Networks, Synced Review, accessed February 18, 2026.

[9] New research in Hopfield Networks, GenAI Engineer, accessed February 18, 2026.

[10] Modern Hopfield Networks meet Encoded Neural Representations, arXiv, accessed February 18, 2026.

[11] Analyzing Immunomes using Sequence Embedding and Network Analysis, ResearchGate, accessed February 18, 2026.

[12] Dense Associative Memory for Pattern Recognition, NIPS, accessed February 18, 2026.

[13] Modern Hopfield Networks and Attention for Immune Repertoire Classification, bioRxiv, accessed February 18, 2026.

[14] Implement the new Hopfield Network with PyTorch, Medium, accessed February 18, 2026.

[15] Modern Hopfield Networks and Attention for Immune Repertoire Classification, arXiv, accessed February 18, 2026, https://arxiv.org/abs/2007.13505

[16] ml-jku/hopfield-layers, GitHub, accessed February 18, 2026.

[17] Conv-CoA: Improving Open-domain Question Answering in Large Language Models via Conversational Chain-of-Action, arXiv, accessed February 18, 2026.

[18] AdvEvo-MARL: Shaping Internalized Safety through Adversarial Co-Evolution, ResearchGate, accessed February 18, 2026.

[19] An implementation of the Hopfield Network using PyTorch, GitHub, accessed February 18, 2026.

[20] Bio-Inspired Neural Computing, GitHub Gist, accessed February 18, 2026.

[21] SQuAD Question Answering using Multi-Perspective Matching, accessed February 18, 2026.

[22] Associative Transformer, arXiv, accessed February 18, 2026.

[23] HLTRI at W-NUT 2020 Shared Task-3: COVID-19 Event Extraction, ACL Anthology, accessed February 18, 2026.

[24] A Fast and High Quality Multilevel Scheme for Partitioning Irregular Graphs, SIAM Journal on Scientific Computing, accessed February 18, 2026.

[25] Improved Robustness and Hyperparameter Selection in Modern Hopfield Networks, accessed February 18, 2026.

[26] Transformers as Intrinsic Optimizers: Forward Inference through the Energy Principle, arXiv, accessed February 18, 2026.

[27] Energy Transformer, NeurIPS, accessed February 18, 2026.

[28] On Sparse Modern Hopfield Model, NIPS, accessed February 18, 2026.

[29] STanHop: Sparse Tandem Hopfield Model for Memory-Enhanced Time Series Prediction, accessed February 18, 2026.

[30] Hideya Ochiai's research works, ResearchGate, accessed February 18, 2026.

[31] Scalable Memory QA with Multi-Signal Associative Retrieval, arXiv, accessed February 18, 2026.

[32] AI Meets Brain: A Unified Survey on Memory Systems, arXiv, accessed February 18, 2026.

[33] Modern Hopfield Networks with Continuous-Time Memories, arXiv, accessed February 18, 2026.

[34] Paper Digest: Recent Papers on Transformer, accessed February 18, 2026.

[35] On the Role of Hidden States of Modern Hopfield Network in Transformer, arXiv, accessed February 18, 2026.

[36] Leveraging Modern Hopfield Networks in Large Language Models, accessed February 18, 2026.

[37] In search of AI algorithms that mimic the brain, IBM Research, accessed February 18, 2026.

TABLE I
SYSTEMATIC COMPENDIUM OF ASSOCIATIVE MEMORY LITERATURE

| Year | Paper Title & Authors | | Proposed Methodology / Framework | QoS Parameters / Problems Solved | Limitations |
|---|---|---|---|---|---|
| 2016 | Dense Associative Memory for Pattern Recognition | (Krotov & Hopfield) | Introduces continuous variables and higher-order polynomial interactions to replace binary Hebbian weights. | Solves the $0.138N$ linear scaling limit of classical networks. Achieves super-linear storage. | Computationally heavy; requires rigid activation functions to maintain global stability on the energy landscape. |
| 2020 | Hopfield Networks is All You Need | (Ramsauer et al.) | Proposes the continuous Modern Hopfield Network utilizing the LogSumExp energy function. Proves mathematical equivalence to Transformer attention. | Solves retrieval bottlenecks. Grants exponential storage capacity $2^{d/2}$. Global convergence typically occurs in a single step. | Requires retaining all raw memory patterns explicitly in the dense matrix $X$, leading to high memory footprint. |
| 2020 | Modern Hopfield Networks and Attention for Immune Repertoire Classification | (Widrich et al.) | Introduces the DeepRC framework, integrating transformer-like attention (MHN) into deep learning for massive Multiple Instance Learning. | Solves massive MIL problems characterized by extremely low witness rates and unprecedented numbers of instances. | Hyperparameter sensitivity to the inverse temperature; model evaluated specifically on biological sequence data. |
| 2022 | Modern Hopfield Networks for graph embedding | (Liang et al.) | MHN-based network embedding. Learns associations between node content and neighbors using recurrent dynamics to recover masked states. | Achieves competitive QoS in downstream tasks like node classification, link prediction, and graph coarsening against baselines. | Struggles to adapt to rapidly dynamic graphs; requires offline calculation of static topological neighbor matrices. |
| 2023 | Energy Transformer | (Hoover et al.) | The Energy Transformer (ET) architecture replaces standard interleaved transformer blocks with an energy-minimizing sequence of attention layers. | Improves theoretical grounding of token relationships. Yields coherent pattern completion by minimizing global energy. | Architecture requires rigid adherence to specific energy functions, limiting general adaptability to off-the-shelf QA tasks. |
| 2023 | Sparse modern hopfield networks | (Hu et al.) | Introduces the $\alpha$-entmax family of sparse probability mappings to regularize energy functionals | Solves the emergence of spurious states and metastability. Enhances pattern separation by mapping to | Gradient propagation through non-differentiable sparse mapping functions can be highly unstable |